

Co-processor instruction set

The vector processor will be a co-processor to a RISC-V core.

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit	Comment
R-type	funct7	rs2	rs1	funct3	rd	opcode	Register-register ALU
I-type	imm		rs1	funct3	rd	opcode	Register-immediate ALU, load
S-type	imm	rs2	rs1	funct3	imm	opcode	Store
U-type	imm				rd	opcode	Loading immediates

- **opcode** - general class of operation e.g. OP-IMM for register-immediate ALU operations
- **rs1** - first source register
- **rs2** - second register
- **rd** - destination register
- **imm** - immediate value
- **funct3** - select type of operation e.g. sub, mul or add
- **funct7** - spill over details of operation

RISC-V ISAs have prefixes e.g. the standard integer instruction set has the '11' prefix and thus has 30-bits of instruction space. Thus there are 3 other available 30 bit instruction spaces available although they can be occupied by the compressed-instruction extensions. As BRAM is relatively plentiful and can easily provide a 32-bit instruction read port, we are not going to implement compressed instructions and instead implement the vector processor as a full 30 bit extension with a prefix of '10'.

VSOP

Perform arithmetic operations between vectors. As we experiment with different numerics, quite how these operations are performed will change e.g. saturated or stochastic rounded but the ISA will remain the same

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
R-type	funct7	rs2	rs1	funct3	rd	0000010

- **rd** selects target vector register
- **rs1** and **rs2** select source vector registers
- **funct3** selects
 - VADD (000): $rd[i] = \text{sat}(rs1[i] + rs2[i])$
 - VSUB (010): $rd[i] = \text{sat}(rs1[i] - rs2[i])$
 - VAND (011): $rd[i] = rs1[i] \& rs2[i]$
 - VSL (001): $rd[i] = rs1[i] \ll rs2[i]$
 - VSRA (101): $rd[i] = rs1[i] \gg rs2[i]$
 - VMUL (100): $rd[i] = \text{round}(rs1[i] * rs2[i]) \gg \text{shift}$
- **funct7**:

Field-size	1 bit	2 bit	4 bit
funct7	saturate	rounding	shift

- **saturate** when performing an addition or subtraction (funct3 is 000 or 010), if this bit is set saturate the result between -32768 and 32767
- **shift** when performing a multiplication (funct3 is VMUL), the result is 32-bit so we need to select which 16-bit component we actually want using a barrel shifter
- **round** when performing a multiplication (funct3 is 100), selects rounding mode:
 - 00 - standard round to zero $\text{round}(x) = x$
 - 01 - round to nearest $\text{round}(x) = x + (1 \ll (\text{shift} - 1))$
 - 10 - stochastic $\text{round}(x) = x + (\text{RNG} \& ((1 \ll \text{shift}) - 1))$

VSOP-IMM

Perform arithmetic operations between vectors and immediates. These are mostly shift operations as you can't fit a 16-bit immediate into a 2 operand RISC-V instruction:

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
I-type	imm	rs1	funct3	rd	0100101	

- **rd** selects target vector register
- **rs1** elect source vector register
- **funct3** selects
 - VSLLI (000): $\text{rd}[i] = \text{rs1}[i] \ll \text{shift}$
 - VSRAI (001): $\text{rd}[i] = \text{round}(\text{rs1}[i]) \gg \text{shift}$
- **imm:**

Field-size	6 bit	2 bit	4 bit
imm	000000	rounding	shift

- **shift** select shift to apply
- **round** when performing a right shift (funct3 is 001), selects rounding mode:
 - 00 - standard round to zero $\text{round}(x) = x$
 - 01 - round to nearest $\text{round}(x) = x + (1 \ll (\text{shift} - 1))$
 - 10 - stochastic $\text{round}(x) = x + (\text{RNG} \& ((1 \ll \text{shift}) - 1))$

VTST

Compare operands writes resultant 32-bit masks to scalar register file.

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
R-type	0000000	rs2	rs1	funct3	rd	0001010

- **rd** selects target scalar register
- **rs1** and **rs2** select source vector registers

- **funct3** selects type of comparison
 - VTEQ (000): $rd = ((rs1[i] == rs2[i]) \ll i)$
 - VTNE (010): $rd = ((rs1[i] != rs2[i]) \ll i)$
 - VTLT (100): $rd = ((rs1[i] < rs2[i]) \ll i)$
 - VTGE (110): $rd = ((rs1[i] >= rs2[i]) \ll i)$

VSEL

Ternary operator. Used for implementing non-uniform control flow.

$rd[i] = (rs1 \& (1 \ll i)) ? rs2[i] : rd[i]$

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
R-type	0000000	rs2	rs1	000	rd	0001110

- **rd** selects target vector register
- **rs1** selects scalar register containing mask
- **rs2** selects vector register to move if mask is true

VLUI

This isn't strictly necessary as you could just load immediates into scalar registers and move them into vector **but**, there may be more register pressure in the scalar register file (especially if we want to respect the RISC-V calling convention):

$rd[i] = imm$

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
U-type	imm				rd	0000110

- **imm:**

Field-size	4 bit	16 bit
imm	zero	Immediate value

- **rd** selects vector register to load into

VFILL

Fills a vector register with the lower 16-bit of scalar register value. Loading directly from scalar memory into vector register would be possible but would involve integrating with XIF memory interface.

$rd[i] = (rs1 \& 0x80000000) \gg 16 \mid (rs1 \& 0x7FFF)$

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
I-type	0000000000000		rs1	000	rd	0011010

- **rd** selects target vector register
- **rs1** selects scalar register containing value to fill. This value is treated as signed so sign bit needs to be combined with

VEXTRACT

Extract the value from one lane of vector register, sign extend to 32-bits and write to scalar register.

$rd = (rs1[imm] \ll 16) \gg 16$

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit
I-type	0000000	imm	rs1	001	rd	0011010

- **rd** selects target scalar register
- **rs1** selects vector register value to extract. This value is treated as signed so sign bit needs to be combined with
- **imm** selects lane to extract

VSPC

Special functions used for random number generation etc. Random number generator is Xoroshiro32++. Because we assume vector lane values are signed, shifts down so you get random numbers in (0, 1) in s0.15 fixed point.

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit	Comment
I-type	0000000000000		rs1	funct3	rd	0100010	Special

- **rd** selects vector target register
- **rs1** source vector register if appropriate
- **funct3** selects type of special function
 - VRNG (000): $rd = rng() \gg 1$

VLOAD

Reads a vector register from vector or lane local memory. VLOADR0 and VLOADR1 are special versions used to load RNG seed into special seed registers.

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit	Comment
I-type	imm		rs1	000	rd	0010010	Vector load

- **imm** selects offset in bytes from base address for vector BRAM, this must be 64-byte vector-size aligned)
- **rd** selects vector register to load into
- **rs1** selects register with base address in bytes. For local loads this is a vector register and for vector loads this is a scalar register. Local loads must be half-word aligned and vector loads must be vector (64 byte) aligned.
- **funct3**
 - VLOADV (000): $rd[i] = VMEM[(rs1 + imm) / 2] + i]$
 - VLOADL (010): $rd[i] = VLOCAL_i[(rs1[i] + imm) / 2]$
 - VLOADR0 (001): $rng[0] = VMEM[(rs1 + imm) / 2] + i]$
 - VLOADR1 (101): $rng[1] = VMEM[(rs1 + imm) / 2] + i]$

VSTORE

Stores the contents of a vector register to vector or lane local memory:

$rs1[(imm / 2) + i] = rs2[i]$

Field-size	7 bit	5 bit	5 bit	3 bit	5 bit	7 bit	Comment
S-type	imm	rs2	rs1	funct3	imm	0010110	Store

- **imm** selects offset in bytes from base address (for vector BRAM, this must be 64-byte vector-size aligned)
- **rs1** selects register with base address in bytes. For local stores this is a vector register and for vector stores this is a scalar register. Local stores must be half-word aligned and vector stores must be vector (64 byte) aligned. For local loads, write enable if $rs1[i] \geq 0$
- **rs2** selects vector register to store
- **funct3**
 - VSTOREV (000): $VMEM[(rs1 + imm) / 2] + i] = rs2[i]$
 - VSTOREL (010): $rd[i] = VLOCAL_i[(rs1[i] + imm) / 2]$

CSRs

FeNN's DMA controller can be controlled by interacting with a number of registers, exposed as CSRS using the standard RISC-V CSR instruction set. MM2S refers to copies from the physical addresses in the system memory map (typically a userspace DMA buffer configured from the PS) to FeNN's vector memory. S2MM refers to copies from FeNN's vector memory to physical addresses in the system memory map.

Name	CSR	Access	Description
MM2S_SRC_ADDR	BC0	RW	32 bit source physical address in system memory map to copy from. The address must be 64 byte aligned.
MM2S_DST_ADDR	BC1	RW	32 bit destination address in vector memory to copy to. The address must be 64 byte aligned.

MM2S_COUNT	BC2	RW	Number of bytes to copy from memory-mapped to vector memory. The count must be a multiple of 64.	
MM2S_STATUS	FC0	R	Bit	Purpose
			0	Unused
			1	Is S2MM FSM idling (debug use)
			2	Is S2MM FSM starting (debug use)
			3	Is S2MM FSM transferring (debug use)
			4	Slave error occurred on last transfer
			5	Decode error occurred on last transfer
			6	Internal error occurred on last transfer
			7	Last transfer completed successfully
MM2S_CONTROL	BC3	W	Writing a 1 to lowest bit of control register starts the MM2S transfer described by MM2S_SRC_ADDR, MM2S_DST_ADDR and MM2S_COUNT registers	
S2MM_SRC_ADDR	BD0	RW	32 bit source address in vector memory to copy from. The address must be 64 byte aligned.	
S2MM_DST_ADDR	BD1	RW	32 bit destination physical address in system memory map to copy to. The address must be 64 byte aligned.	
S2MM_COUNT	BD2	RW	Number of bytes to copy from vector memory to memory-mapped. The count must be a multiple of 64.	
S2MM_STATUS	FD0	R	Bit	Purpose
			0	Unused
			1	Is S2MM FSM idling (debug use)
			2	Is S2MM FSM starting (debug use)
			3	Is S2MM FSM transferring (debug use)
			4	Slave error occurred on last transfer
			5	Decode error occurred on last transfer
			6	Internal error occurred on last transfer
			7	Last transfer completed successfully

S2MM_CONTROL	BD3	W	Writing a 1 to lowest bit of control register starts the MM2S transfer described by MM2S_SRC_ADDR, MM2S_DST_ADDR and MM2S_COUNT registers
--------------	-----	---	---